

Experiment No. 3
Apply Decision Tree Algorithm on Adult Census Income Dataset and analyze the performance of the model
Date of Performance:
Date of Submission:

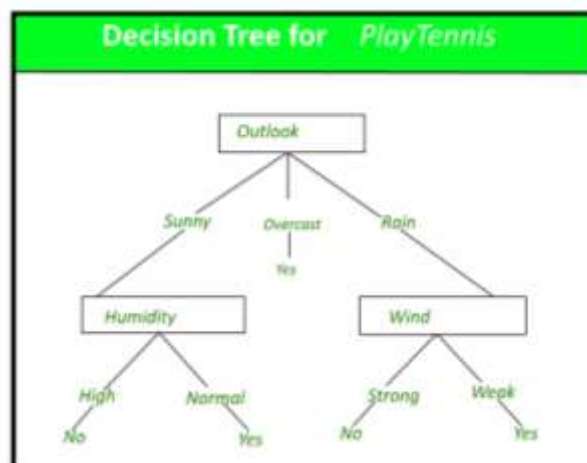


Aim: Apply Decision Tree Algorithm on Adult Census Income Dataset and analyze the performance of the model.

Objective: To perform various feature engineering tasks, apply Decision Tree Algorithm on the given dataset and maximize the accuracy, Precision, Recall, F1 score. Improve the performance by performing different data engineering and feature engineering tasks.

Theory:

Decision Tree is the most powerful and popular tool for classification and prediction. A Decision tree is a flowchart-like tree structure, where each internal node denotes a test on an attribute, each branch represents an outcome of the test, and each leaf node (terminal node) holds a class label.



Dataset:

Predict whether income exceeds \$50K/yr based on census data. Also known as "Adult" dataset.

Attribute Information:

Listing of attributes:

>50K, <=50K.



Vidyavardhini's College of Engineering & Technology

Department of Computer Engineering

age: continuous.

workclass: Private, Self-emp-not-inc, Self-emp-inc, Federal-gov, Local-gov, State-gov, Without-pay, Never-worked.

fnlwgt: continuous.

education: Bachelors, Some-college, 11th, HS-grad, Prof-school, Assoc-acdm, Assoc-voc, 9th, 7th-8th, 12th, Masters, 1st-4th, 10th, Doctorate, 5th-6th, Preschool.

education-num: continuous.

marital-status: Married-civ-spouse, Divorced, Never-married, Separated, Widowed, Married-spouse-absent, Married-AF-spouse.

occupation: Tech-support, Craft-repair, Other-service, Sales, Exec-managerial, Prof-specialty, Handlers-cleaners, Machine-op- Inspct, Adm-clerical, Farming-fishing, Transport-moving, Priv-house-serv, Protective-serv, Armed-Forces.

relationship: Wife, Own-child, Husband, Not-in-family, Other-relative, Unmarried.

race: White, Asian-Pac-Islander, Amer-Indian-Eskimo, Other, Black.

sex: Female, Male.

capital-gain: continuous.

capital-loss: continuous.

hours-per-week: continuous.

native-country: United-States, Cambodia, England, Puerto-Rico, Canada, Germany, Outlying-US(Guam-USVI-etc), India, Japan, Greece, South, China, Cuba, Iran, Honduras, Philippines, Italy, Poland, Jamaica, Vietnam, Mexico, Portugal, Ireland, France, Dominican-Republic,



Laos, Ecuador, Taiwan, Haiti, Columbia, Hungary, Guatemala, Nicaragua, Scotland, Thailand, Yugoslavia, El-Salvador, Trinidad&Tobago, Peru, Hong, Holand-Netherlands.

Code:

Conclusion:

1. Discuss about the how categorical attributes have been dealt with during data pre-processing.

The given dataset contained categorical attributes such as workclass, education, marital status, relationship, race, sex, native country and income these categorical attributes haven been dealt using data preprocessing techniques such as label encoding. Label encoding is a technique used in machine learning and data analysis to convert categorical variables into numerical format.

2. Discuss the hyper-parameter tuning done based on the decision tree obtained.

Hyperparameter tuning, also known as hyperparameter optimization, is the process of finding the best set of hyperparameters for a machine learning model. Hyperparameters are settings or configurations that are not learned from the data but are set prior to training a model.

3. Comment on the accuracy, confusion matrix, precision, recall and F1 score obtained.

Accuracy: Model accuracy is about 81%. This means the model correctly predicted the class label.

Precision: Precision measures the accuracy of the positive predictions and the precision score obtained by our model is 59%

Recall: The recall rate is about 62%. This means that the model can only correctly identify about 62% of all real cases.

F1 Score: The F1 score of 60.8% is between precision and recall, providing a balanced view of model performance.



Vidyavardhini's College of Engineering & Technology

Department of Computer Engineering

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn import linear_model
from sklearn.tree import DecisionTreeClassifier
df= pd.read_csv('adult.csv')
```

```
df= pd.read_csv('adult.csv')
df.head()
```

	age	workclass	fnlwgt	education	education.num	marital.status	occupation	relationship	race	
0	90	?	77053	HS-grad	9	Widowed	?	Not-in-family	White	Fe
1	82	Private	132870	HS-grad	9	Widowed	Exec-managerial	Not-in-family	White	Fe
2	66	?	186061	Some-college	10	Widowed	?	Unmarried	Black	Fe
3	54	Private	140359	7th-8th	4	Divorced	Machine-op-inspct	Unmarried	White	Fe
4	41	Private	264663	Some-college	10	Separated	Prof-specialty	Own-child	White	Fe

```
df.shape
(32561, 15)
```

```
df.describe()
```

	age	fnlwgt	education.num	capital.gain	capital.loss	hours.p
count	32561.000000	3.256100e+04	32561.000000	32561.000000	32561.000000	32561.000000
mean	38.581647	1.897784e+05	10.080679	1077.648844	87.303830	40.620855
std	13.640433	1.055500e+05	2.572720	7385.292085	402.960219	11.516961
min	17.000000	1.228500e+04	1.000000	0.000000	0.000000	0.000000
25%	28.000000	1.178270e+05	9.000000	0.000000	0.000000	40.000000
50%	37.000000	1.783560e+05	10.000000	0.000000	0.000000	40.000000
75%	48.000000	2.370510e+05	12.000000	0.000000	0.000000	40.000000
max	90.000000	1.481705e+06	16.000000	99999.000000	1356.000000	99.000000

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 32561 entries, 0 to 32560
Data columns (total 15 columns):
#   Column                Non-Null Count  Dtype
---  -
0   age                    32561 non-null  int64
1   workclass              32561 non-null  object
2   fnlwgt                 32561 non-null  int64
3   education              32561 non-null  object
4   education.num          32561 non-null  int64
5   marital.status         32561 non-null  object
6   occupation              32561 non-null  object
7   relationship           32561 non-null  object
```

```

8   race           32561 non-null object
9   sex            32561 non-null object
10  capital.gain    32561 non-null int64
11  capital.loss    32561 non-null int64
12  hours.per.week  32561 non-null int64
13  native.country  32561 non-null object
14  income          32561 non-null object
dtypes: int64(6), object(9)
memory usage: 3.7+ MB

```

```
df['income'].value_counts()
```

```

<=50K    24720
>50K      7841
Name: income, dtype: int64

```

```
df['sex'].value_counts()
```

```

Male      21790
Female    10771
Name: sex, dtype: int64

```

```
df['native.country'].value_counts()
```

```

United-States    29170
Mexico           643
?                583
Philippines      198
Germany          137
Canada           121
Puerto-Rico     114
El-Salvador      106
India            100
Cuba             95
England          90
Jamaica          81
South            80
China            75
Italy            73
Dominican-Republic 70
Vietnam          67
Guatemala        64
Japan            62
Poland           60
Columbia         59
Taiwan           51
Haiti            44
Iran             43
Portugal         37
Nicaragua        34
Peru             31
Greece           29
France           29
Ecuador          28
Ireland          24
Hong             20
Cambodia         19
Trinidad&Tobago  19
Laos             18
Thailand         18
Yugoslavia       16
Outlying-US(Guam-USVI-etc) 14
Hungary          13
Honduras         13
Scotland         12
Holand-Netherlands 1
Name: native.country, dtype: int64

```

```
df['workclass'].value_counts()

Private          22696
Self-emp-not-inc  2541
Local-gov        2093
?                1836
State-gov        1298
Self-emp-inc     1116
Federal-gov      960
Without-pay      14
Never-worked     7
Name: workclass, dtype: int64
```

```
df['occupation'].value_counts()

Prof-specialty    4140
Craft-repair      4099
Exec-managerial   4066
Adm-clerical      3770
Sales             3650
Other-service     3295
Machine-op-inspct 2002
?                1843
Transport-moving  1597
Handlers-cleaners 1370
Farming-fishing   994
Tech-support      928
Protective-serv   649
Priv-house-serv   149
Armed-Forces      9
Name: occupation, dtype: int64
```

```
df.replace('?', np.NaN,inplace = True)
df.head()
```

tal.status	occupation	relationship	race	sex	capital.gain	capital.loss	hours.per.week	native.cc
Widowed	NaN	Not-in-family	White	Female	0	4356	40	United-
Widowed	Exec-managerial	Not-in-family	White	Female	0	4356	18	United-
Widowed	NaN	Unmarried	Black	Female	0	4356	40	United-
Divorced	Machine-op-inspct	Unmarried	White	Female	0	3900	40	United-
Separated	Prof-specialty	Own-child	White	Female	0	3900	40	United-

```
#fills missing values (NaNs) with the most recent non-missing value
df.fillna(method = 'ffill', inplace = True)
```

```
from sklearn.preprocessing import LabelEncoder
```

```
le = LabelEncoder()
df['workclass'] = le.fit_transform(df['workclass'])
df['education'] = le.fit_transform(df['education'])
df['marital.status'] = le.fit_transform(df['marital.status'])
df['occupation'] = le.fit_transform(df['occupation'])
df['relationship'] = le.fit_transform(df['relationship'])
df['race'] = le.fit_transform(df['race'])
df['sex'] = le.fit_transform(df['sex'])
df['native.country'] = le.fit_transform(df['native.country'])
df['income'] = le.fit_transform(df['income'])
```



```
df.head()
```

	age	workclass	fnlwgt	education	education.num	marital.status	occupation	relationship	race	sex
0	90	8	77053	11	9	6	14	1	4	0
1	82	3	132870	11	9	6	3	1	4	0
2	66	3	186061	15	10	6	3	4	2	0
3	54	3	140359	5	4	0	6	4	4	0
4	41	3	264663	15	10	5	9	3	4	0

```
df.describe()
```

	onship	race	sex	capital.gain	capital.loss	hours.per.week	native.
000000	32561.000000	32561.000000	32561.000000	32561.000000	32561.000000	32561.000000	3256
446362	3.665858	0.669205	1077.648844	87.303830	40.437456	3	
606771	0.848806	0.470506	7385.292085	402.960219	12.347429		
000000	0.000000	0.000000	0.000000	0.000000	1.000000		
000000	4.000000	0.000000	0.000000	0.000000	40.000000	3	
000000	4.000000	1.000000	0.000000	0.000000	40.000000	3	
000000	4.000000	1.000000	0.000000	0.000000	45.000000	3	
000000	4.000000	1.000000	99999.000000	4356.000000	99.000000	4	

```
#Splitting the data set into features and outcome
X = df.drop(['income'], axis=1)
Y = df['income']
```

```
df.isnull().sum()
```

age	0
workclass	0
fnlwgt	0
education	0
education.num	0
marital.status	0
occupation	0
relationship	0
race	0
sex	0
capital.gain	0
capital.loss	0
hours.per.week	0
native.country	0
income	0

dtype: int64

```
df.duplicated().sum()
```

24

```
df=df.dropna()
```

```
#Splitting the data into test data and training data

from sklearn.model_selection import train_test_split

X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size = 0.2, random_state = 42)

X_train.head()
```

	age	workclass	fnlwgt	education	education.num	marital.status	occupation
5514	26	3	256263	11	9	4	2
19777	24	3	170277	11	9	4	7
10781	36	3	75826	9	13	0	0
32240	22	6	24395	15	10	2	0
9876	31	1	356689	9	13	2	9

```
Y_train = Y_train.replace((np.inf, -np.inf, np.nan), 0).reset_index(drop=True)
```

```
Y_train.head()
```

```
0    0
1    0
2    0
3    0
4    0
Name: income, dtype: int64
```

```
Y_test = Y_test.replace((np.inf, -np.inf, np.nan), 0).reset_index(drop=True)
```

```
from sklearn.tree import DecisionTreeClassifier
dec_tree = DecisionTreeClassifier(random_state=42)
dec_tree.fit(X_train, Y_train)
Y_pred_dec_tree = dec_tree.predict(X_test)
```

```
from sklearn.metrics import accuracy_score
from sklearn.metrics import f1_score
from sklearn.metrics import precision_score
from sklearn.metrics import recall_score
```

```
print('Decision Tree Classifier:')
print('Accuracy score:',accuracy_score(Y_test, Y_pred_dec_tree) * 100)
print('Precision :',precision_score(Y_test,Y_pred_dec_tree) *100)
print('Recall: ',recall_score(Y_test,Y_pred_dec_tree)* 100)
print('F1 score: ',f1_score(Y_test,Y_pred_dec_tree) *100)
```

```
Decision Tree Classifier:
Accuracy score: 81.1761093198219
Precision : 59.749216300940446
Recall: 62.0039037085231
F1 score: 60.85568326947638
```

✓ 0s completed at 12:24 AM

● ×