



Experiment No. 7
Apply Dimensionality Reduction on Adult Census Income Dataset and analyze the performance of the model
Date of Performance:
Date of Submission:



Aim: Apply Dimensionality Reduction on Adult Census Income Dataset and analyze the performance of the model.

Objective: Able to perform various feature engineering tasks, perform dimensionality reduction on the given dataset and maximize the accuracy, Precision, Recall, F1 score.

Theory:

In machine learning classification problems, there are often too many factors on the basis of which the final classification is done. These factors are basically variables called features. The higher the number of features, the harder it gets to visualize the training set and then work on it. Sometimes, most of these features are correlated, and hence redundant. This is where dimensionality reduction algorithms come into play. Dimensionality reduction is the process of reducing the number of random variables under consideration, by obtaining a set of principal variables. It can be divided into feature selection and feature extraction.

Dataset:

Predict whether income exceeds \$50K/yr based on census data. Also known as "Adult" dataset.

Attribute Information:

Listing of attributes:

>50K, <=50K.

age: continuous.

workclass: Private, Self-emp-not-inc, Self-emp-inc, Federal-gov, Local-gov, State-gov,

Without-pay, Never-worked. fnlwgt: continuous.



Vidyavardhini's College of Engineering & Technology

Department of Computer Engineering

education: Bachelors, Some-college, 11th, HS-grad, Prof-school, Assoc-acdm, Assoc-voc, 9th, 7th-8th, 12th, Masters, 1st-4th, 10th, Doctorate, 5th-6th, Preschool. education-num: continuous.

marital-status: Married-civ-spouse, Divorced, Never-married, Separated, Widowed, Marriedspouse-absent, Married-AF-spouse.

occupation: Tech-support, Craft-repair, Other-service, Sales, Exec-managerial, Prof-specialty, Handlers-cleaners, Machine-op-ops, Adm-clerical, Farming-fishing, Transport-moving, Priv-house-serv, Protective-serv, Armed-Forces.

relationship: Wife, Own-child, Husband, Not-in-family, Other-relative, Unmarried.

race: White, Asian-Pac-Islander, Amer-Indian-Eskimo, Other, Black. sex: Female,

Male. capital-gain: continuous. capital-loss: continuous. hours-per-week:

continuous.

native-country: United-States, Cambodia, England, Puerto-Rico, Canada, Germany, Outlying-US(Guam-USVI-etc), India, Japan, Greece, South, China, Cuba, Iran, Honduras, Philippines, Italy, Poland, Jamaica, Vietnam, Mexico, Portugal, Ireland, France, DominicanRepublic, Laos, Ecuador, Taiwan, Haiti, Columbia, Hungary, Guatemala, Nicaragua, Scotland, Thailand, Yugoslavia, El-Salvador, Trinidad & Tobago, Peru, Hong, HolandNetherlands.

Code:



Conclusion:

1. Comment on the impact of dimensionality reduction on the accuracy, precision, recall and F1Score

Dimensionality reduction with PCA has a modest impact on performance.

The accuracy drops slightly from 84.9% to 82.5%.

Precision for '>50K' decreases from 0.73 to 0.72, while precision for '<=50K' remains at 0.87.

Recall for '>50K' decreases from 0.61 to 0.54, but recall for '<=50K' stays high at 0.94.

F1-scores show a decline, with '>50K' dropping from 0.66 to 0.61, and '<=50K' remaining high at 0.90.

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
```

```
data = pd.read_csv('adult.csv')
```

```
print(data)
```

```
0      90  Private  77053  HS-grad  9  Widowed
1      82  Private 132870  HS-grad  9  Widowed
2      66      ? 186061  Some-college 10  Widowed
3      54  Private 140359  7th-8th  4  Divorced
4      41  Private 264663  Some-college 10  Separated
...
32556  22  Private 310152  Some-college 10  Never-married
32557  27  Private 257302  Assoc-acdm 12  Married-civ-spouse
32558  40  Private 154374  HS-grad  9  Married-civ-spouse
32559  58  Private 151910  HS-grad  9  Widowed
32560  22  Private 201490  HS-grad  9  Never-married

0      ?  Not-in-family  White  Female  0
1  Exec-managerial  Not-in-family  White  Female  0
2      ?  Unmarried  Black  Female  0
3  Machine-op-inspct  Unmarried  White  Female  0
4  Prof-specialty  Own-child  White  Female  0
...
32556  Protective-serv  Not-in-family  White  Male  0
32557  Tech-support  Wife  White  Female  0
32558  Machine-op-inspct  Husband  White  Male  0
32559  Adm-clerical  Unmarried  White  Female  0
32560  Adm-clerical  Own-child  White  Male  0

0  4356  40  United-States  <=50K
1  4356  18  United-States  <=50K
2  4356  40  United-States  <=50K
3  3900  40  United-States  <=50K
4  3900  40  United-States  <=50K
...
32556  0  40  United-States  <=50K
32557  0  38  United-States  <=50K
32558  0  40  United-States  >50K
32559  0  40  United-States  <=50K
32560  0  20  United-States  <=50K
```

```
[32561 rows x 15 columns]
```

```
data.describe()
```

	age	fnlwgt	education.num	capital.gain	capital.loss	hours.per.wk
count	32561.000000	3.256100e+04	32561.000000	32561.000000	32561.000000	32561.000000
mean	38.581647	1.897784e+05	10.080679	1077.648844	87.303830	40.437081
std	13.640433	1.055500e+05	2.572720	7385.292085	402.960219	12.347602
min	17.000000	1.228500e+04	1.000000	0.000000	0.000000	1.000000
25%	28.000000	1.178270e+05	9.000000	0.000000	0.000000	40.000000
50%	37.000000	1.783560e+05	10.000000	0.000000	0.000000	40.000000
75%	48.000000	2.370510e+05	12.000000	0.000000	0.000000	45.000000
max	90.000000	1.484705e+06	16.000000	99999.000000	4356.000000	99.000000

```
data.isnull().sum()
```

```
age      0
workclass 0
fnlwgt    0
education 0
education.num 0
```

```
marital.status    0
occupation        0
relationship      0
race              0
sex               0
capital.gain      0
capital.loss      0
hours.per.week    0
native.country    0
income            0
dtype: int64
```

```
import matplotlib.pyplot as mp
import pandas as pd
import seaborn as sb
print(data.corr())

# plotting correlation heatmap
dataplot = sb.heatmap(data.corr(), cmap="YlGnBu", annot=True)

# displaying heatmap
mp.show()
```

```
# Encode categorical features using Label Encoding
categorical_features = ['workclass', 'education', 'marital.status', 'occupation', 'relationship', 'race', 'sex', 'native.country', 'income']
for feature in categorical_features:
    label_encoder = LabelEncoder()
```

c

```
<ipython-input-28-b698e0a536da>:4: FutureWarning: The default value of numeric_only in DataFrame.corr is deprecated. In a future versior
print(data.corr())
```

```
<ipython-input-28-b698e0a536da>:7: FutureWarning: The default value of numeric_only in DataFrame.corr is deprecated. In a future versior
dataplot = sb.heatmap(data.corr(), cmap="YlGnBu", annot=True)
```

```

age          1.000000 -0.076646  0.036527  0.077674  0.057775
fnlwgt       -0.076646  1.000000 -0.043195  0.000432 -0.010252
education.num 0.036527 -0.043195  1.000000  0.122630  0.079923
capital.gain  0.077674  0.000432  0.122630  1.000000 -0.031615
capital.loss  0.057775 -0.010252  0.079923 -0.031615  1.000000
hours.per.week 0.068756 -0.018768  0.148123  0.078409  0.054256


```

hours.per.week

```

age          0.068756
fnlwgt       -0.018768
education.num 0.148123
capital.gain  0.078409
capital.loss  0.054256
hours.per.week 1.000000

```



```

data[feature] = label_encoder.fit_transform(data[feature])

# Separate features and target variable
X = data.drop('income', axis=1)
y = data['income']

print(x)
print(y)

   age workclass  fnlwgt  education  education.num  marital.status \
1    82   Private  132870    HS-grad           9      Widowed
3    54   Private  140359    7th-8th           4      Divorced
4    41   Private  264663  Some-college        10      Separated
5    34   Private  216864    HS-grad           9      Divorced
6    38   Private  150601    10th             6      Separated
...   ...   ...   ...   ...   ...   ...
32556  22   Private  310152  Some-college        10  Never-married
32557  27   Private  257302  Assoc-acdm        12  Married-civ-spouse
32558  40   Private  154374    HS-grad           9  Married-civ-spouse
32559  58   Private  151910    HS-grad           9      Widowed
32560  22   Private  201490    HS-grad           9  Never-married

   occupation  relationship  race  sex  capital.gain \
1  Exec-managerial  Not-in-family  White  Female      0
3  Machine-op-inspct  Unmarried  White  Female      0
4  Prof-specialty  Own-child  White  Female      0
5  Other-service  Unmarried  White  Female      0
6  Adm-clerical  Unmarried  White  Male      0
...   ...   ...   ...   ...   ...
32556  Protective-serv  Not-in-family  White  Male      0
32557  Tech-support  Wife  White  Female      0
32558  Machine-op-inspct  Husband  White  Male      0
32559  Adm-clerical  Unmarried  White  Female      0
32560  Adm-clerical  Own-child  White  Male      0

   capital.loss  hours.per.week  native.country
1          4356             18  United-States
3          3900             40  United-States
4          3900             40  United-States
5          3770             45  United-States
6          3770             40  United-States
...   ...   ...   ...
32556          0             40  United-States
32557          0             38  United-States
32558          0             40  United-States
32559          0             40  United-States
32560          0             20  United-States

[30162 rows x 14 columns]
1      0
3      0
4      0
5      0
6      0
..
32556  0
32557  0
32558  1
32559  0
32560  0
Name: income, Length: 30162, dtype: int64

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Standardize numerical features
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)

# Apply PCA for dimensionality reduction
pca = PCA(n_components=10) # Adjust the number of components as needed
X_train_pca = pca.fit_transform(X_train)
X_test_pca = pca.transform(X_test)

# Train a classifier on the original and reduced-dimension data
classifier_original = RandomForestClassifier(random_state=42)
classifier_original.fit(X_train, y_train)
classifier_pca = RandomForestClassifier(random_state=42)
classifier_pca.fit(X_train_pca, y_train)

```


c

```

classifier_pca.fit(X_train_p a, y_train)

# Evaluate model performance
y_pred_original = classifier_original.predict(X_test)
y_pred_pca = classifier_pca.predict(X_test_pca)

accuracy_original = accuracy_score(y_test, y_pred_original)
accuracy_pca = accuracy_score(y_test, y_pred_pca)

print("Accuracy (Original Data):", accuracy_original)
print("Accuracy (PCA Reduced Data):", accuracy_pca)

    Accuracy (Original Data): 0.8491629371788496
    Accuracy (PCA Reduced Data): 0.8252942151500083

from sklearn.metrics import classification_report, confusion_matrix

# Evaluate model performance on the original data
print("Performance on Original Data:")
print("Confusion Matrix:")
confusion_original = confusion_matrix(y_test, y_pred_original)
print(confusion_original)

report_original = classification_report(y_test, y_pred_original)
print("Classification Report:")
print(report_original)

# Evaluate model performance on PCA-reduced data
print("\nPerformance on PCA Reduced Data:")
print("Confusion Matrix:")
confusion_pca = confusion_matrix(y_test, y_pred_pca)
print(confusion_pca)

report_pca = classification_report(y_test, y_pred_pca)
print("Classification Report:")
print(report_pca)

```

```

Performance on Original Data:
Confusion Matrix:
[[4630  346]
 [ 602  935]]
Classification Report:

```

	precision	recall	f1-score	support
0	0.88	0.93	0.91	4976
1	0.73	0.61	0.66	1537
accuracy			0.85	6513
macro avg	0.81	0.77	0.79	6513
weighted avg	0.85	0.85	0.85	6513

```

Performance on PCA Reduced Data:
Confusion Matrix:
[[4655  321]
 [ 712  825]]
Classification Report:

```

	precision	recall	f1-score	support
0	0.87	0.94	0.90	4976
1	0.72	0.54	0.61	1537
accuracy			0.84	6513
macro avg	0.79	0.74	0.76	6513
weighted avg	0.83	0.84	0.83	6513

```

import numpy as np
import matplotlib.pyplot as plt
import itertools

```

```

def plot_confusion_matrix(cm, classes, title, cmap=plt.cm.Blues):
    plt.imshow(cm, interpolation='nearest', cmap=cmap)
    plt.colorbar()

```

c

```

plt.xticks(tick_marks, classes, rotation=45)
plt.yticks(tick_marks, classes)

fmt = 'd'
thresh = cm.max() / 2.
for i, j in itertools.product(range(cm.shape[0]), range(cm.shape[1])):
    plt.text(j, i, format(cm[i, j], fmt),
             horizontalalignment="center",
             color="white" if cm[i, j] > thresh else "black")

plt.ylabel('True label')
plt.xlabel('Predicted label')
plt.tight_layout()

# Plot confusion matrix for the classifier on original data
plt.figure(figsize=(10, 5))
plt.subplot(1, 2, 1)
plot_confusion_matrix(confusion_original, classes=['<=50K', '>50K'], title="Confusion Matrix (Original Data)")

# Plot confusion matrix for the classifier on PCA-reduced data
plt.subplot(1, 2, 2)
plot_confusion_matrix(confusion_pca, classes=['<=50K', '>50K'], title="Confusion Matrix (PCA Reduced Data)")

plt.tight_layout()
plt.show()

```

